Arhitecturi Paralele
Building a Distributed Database

Prof. Florin Pop
As. Drd. Ing. Cristian Chilipirea
cristian.chilipirea@cs.pub.ro

Elemente preluate din cursul Prof. Ciprian Dobre

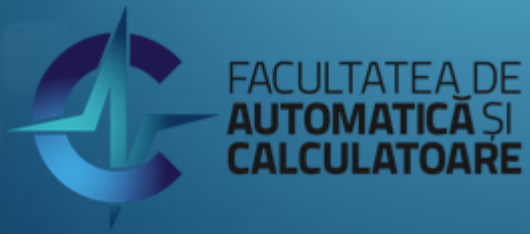


POLITEHNICA
UNIVERSITATEA
BUCUREŞTI
UPB
1818

CAP Theorem

Brewer's Conjecture and the Feasibility of Consistent, Available, Partition-Tolerant Web Services

Seth Gilbert and Nancy Lynch
Laboratory for Computer Science
Massachusetts Institute of Technology
Cambridge, MA 02139
sethg@mit.edu,lynch@theory.lcs.mit.edu

Abstract

When designing distributed web services, there are three pro
desired: consistency, availability, and partition tolerance. It
three. In this note, we prove this conjecture in the asynchr
then discuss solutions to this dilemma in the partially synch

1 Introduction

At PODC 2000, Brewer[1], in an invited talk [2], made the followin
a web service to provide the following three guarantees:

- Consistency
- Availability
- Partition-tolerance


POLITEHNICA
UNIVERSITATEA
BUCUREŞTI
UPB
1818

Bitcoin and Blockchain

Bitcoin: A Peer-to-Peer Electronic Cash System

Satoshi Nakamoto
satoshin@gmx.com
www.bitcoin.org

Abstract. A purely peer-to-peer version of electronic cash would allow online payments to be sent directly from one party to another without going through a financial institution. Digital signatures provide part of the solution, but the main benefits are lost if a trusted third party is still required to prevent double-spending. We propose a solution to the double-spending problem using a peer-to-peer network. The network timestamps transactions by hashing them into an ongoing chain of hash-based proof-of-work, forming a record that cannot be changed without redoing the proof-of-work. The longest chain not only serves as proof of the sequence of events witnessed, but proof that it came from the largest pool of CPU power. As long as a majority of CPU power is controlled by nodes that are not cooperating to attack the network, they'll generate the longest chain and outpace attackers. The network itself requires minimal structure. Messages are broadcast on a best effort basis, and nodes can leave and rejoin the network at will, accepting the longest proof-of-work chain as proof of what happened while they were gone.

1. Introduction

Commerce on the Internet has come to rely almost exclusively on financial institutions serving as trusted third parties to process electronic payments. While the system works well enough for

Cheats on the CAP theorem

Eventual consistency
Availability
Partition-tolerance (**But short**)

Proof of Work

Proof of Work

Pricing via Processing or Combatting Junk Mail

Cynthia Dwork and Moni Naor

IBM Almaden Research Center
650 Harry Road
San Jose, CA 95120

Abstract. We present a computational technique for combatting junk mail in particular and controlling access to a shared resource in general. The main idea is to require a user to compute a moderately hard, but not intractable, function in order to gain access to the resource, thus preventing frivolous use. To this end we suggest several *pricing functions*, based on, respectively, extracting square roots modulo a prime, the Fiat-Shamir signature scheme, and the Ong-Schnorr-Shamir (cracked) signature scheme.

1 Introduction

Recently, one of us returned from a brief vacation, only to find 241 messages in our reader. While junk mail has long been a nuisance in hard (snail) mail, we

Proof of Work

- y = f(x)
  - f() este extrem de lentă computațional
- x = g(y)
  - g() este extrem de rapidă computațional

Proof of Work

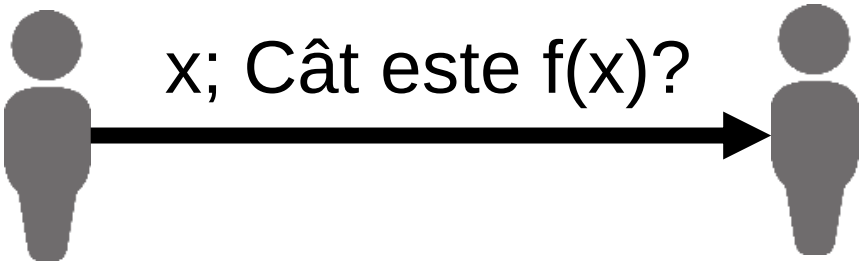

Proof of Work

Processing!!!

Proof of Work

Processing!!!

Proof of Work

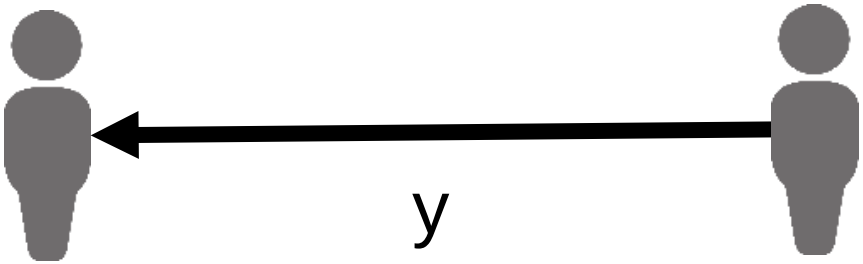

Proof of Work

g(y) == x?
DA - știu că a muncit
NU – încearcă să trișeze

Blockchain

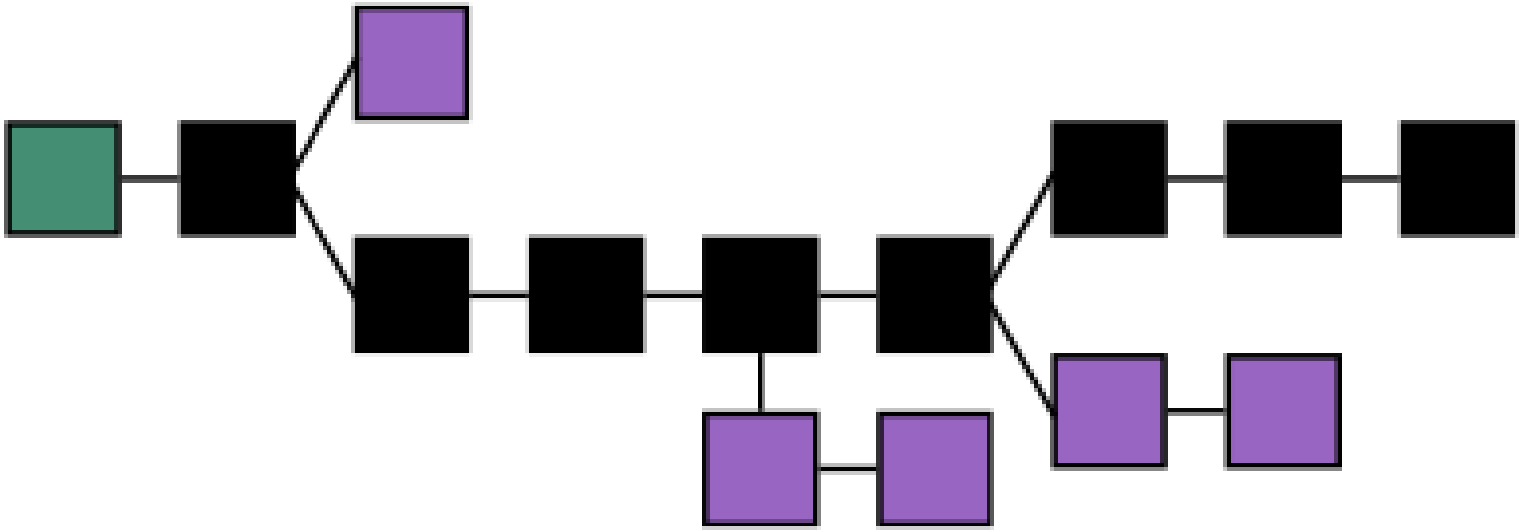

Bitcoin energy consmption

Bitcoin Mining and its Energy Footprint

Karl J. O'Dwyer† and David Malone*

Hamilton Institute
National University of Ireland Maynooth

E-mail: †karl.a.odwyer@nuim.ie *david.malone@nuim.ie

Abstract — **Bitcoin is a digital cryptocurrency that has generated considerable public interest, including both booms in value and busts of exchanges dealing in Bitcoins. One of the fundamental concepts of Bitcoin is that work, called mining, must be done in checking all monetary transactions, which in turn creates Bitcoins as a reward. In this paper we look at the energy consumption of Bitcoin mining. We consider if and when Bitcoin mining has been profitable compared to the energy cost of performing the mining, and conclude that specialist hardware is usually required to make Bitcoin mining profitable. We also show that the power currently used for Bitcoin mining is comparable to Ireland's electricity consumption.**

Bitcoin energy consmption

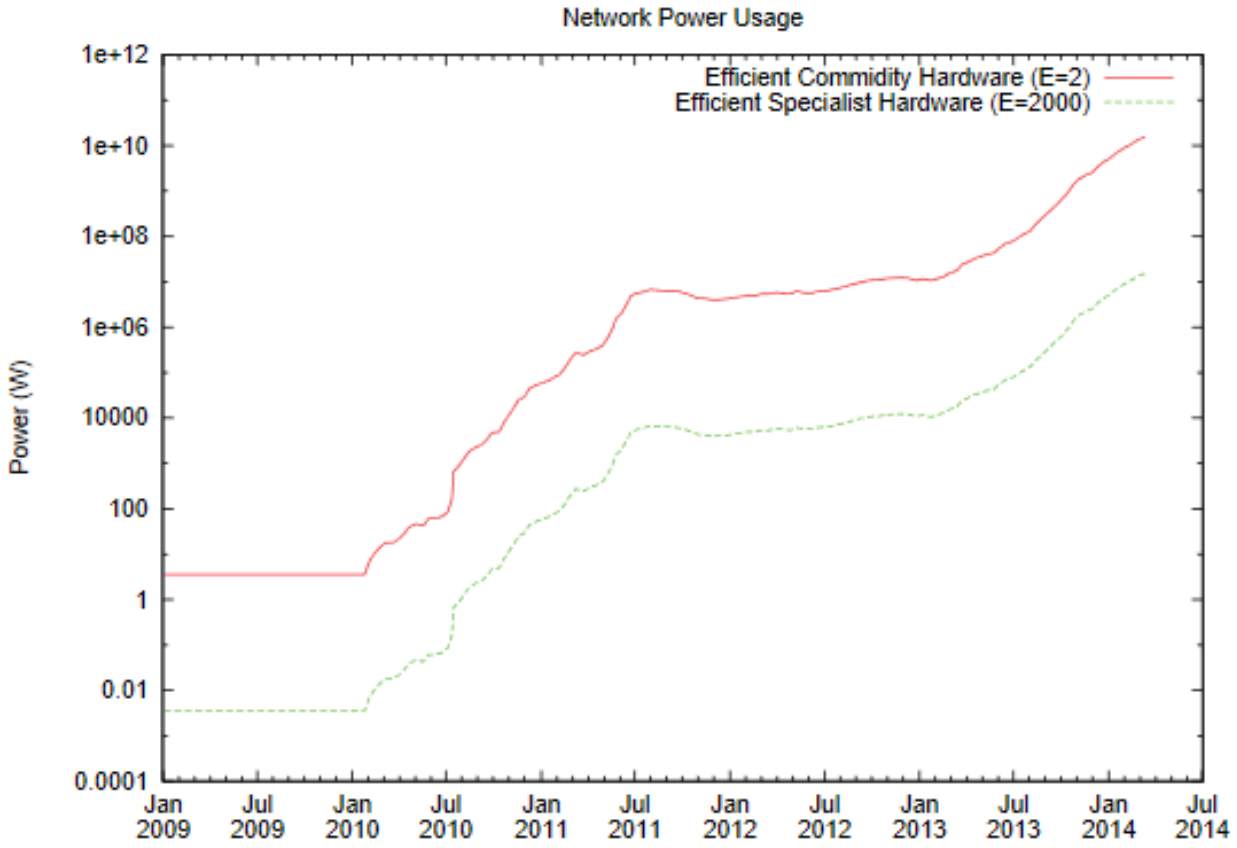

Fig. 5: Estimated Power Consumption of the Bitcoin Mining Network.


POLITEHNICA
UNIVERSITATEA
BUCUREŞTI
UPB
1818

Transactions

ACID

Principles of Transaction-Oriented Database Recovery

THEO HAERDER

Fachbereich Informatik, University of Kaiserslautern, West Germany

ANDREAS REUTER[1]

IBM Research Laboratory, San Jose, California 95193

er, a terminological framework is provided for descri
covery schemes for database systems in a conceptual
ation-dependent way. By introducing the terms mate
n strategy, and checkpoint, we obtain a means for cl
ations from a unified viewpoint. This is complement
logging techniques, which are precisely defined by u
these criteria are related to all relevant questions a

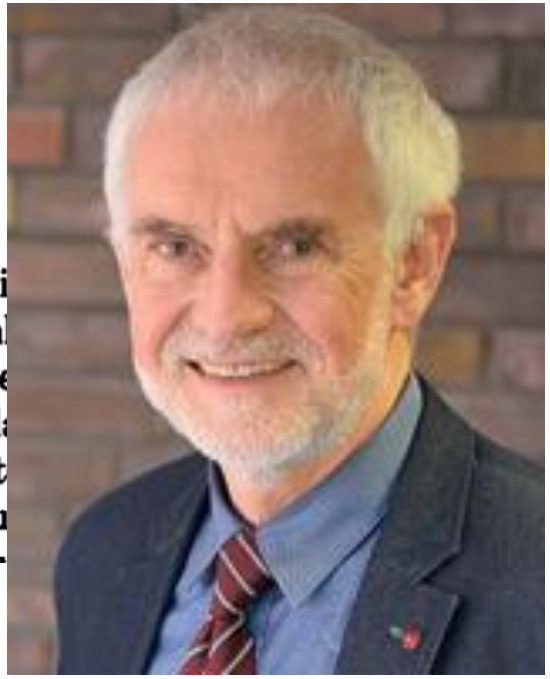


POLITEHNICA
UNIVERSITATEA
BUCUREŞTI
UPB
1818

Distributed Hash Table Chord

Chord: A Scalable Peer-to-peer Lookup Service for Internet Applications

Ion Stoica; Robert Morris, David Karger, M. Frans Kaas
MIT Laboratory for Computer Science
chord@lcs.mit.edu
http://pdos.lcs.mit.edu/chord/

Abstract

A fundamental problem that confronts peer-to-peer applications is to efficiently locate the node that stores a particular data item. This paper presents *Chord*, a distributed lookup protocol that addresses this problem. Chord provides support for just one operation: given a key, it maps the key onto a node. Data location can be easily implemented on top of Chord by associating a key with each data item, and storing the key/data item pair at the node to which the key maps. Chord adapts efficiently as nodes join and leave the system, and can answer queries even if the system is continuously changing. Results from theoretical analysis, simulations, and experiments show that Chord is scalable, with communication cost and the state maintained by each node scaling logarithmically with the number of Chord nodes.

and involves relati
and leave the syste
Previous work o
aware of most othe
scale to large numb
"routing" informati
routing table is dis
communicating wi
an N-node system
$O(\log N)$ other no
sages to other node
nodes join and leav
event results in no
Three features th
peer lookup protoc

Amazon DynamoDB

Caching

Peer to Peer

Geolocation

Content Delivery Networks

Load Balancing

Cloud Computing

Spark?