



# Arhitecturi Paralele

## Consens în Sisteme Distribuite

Lect. Dr. Ing. Cristian Chilipirea  
[cristian.chilipirea@mta.ro](mailto:cristian.chilipirea@mta.ro)



## Dorim sisteme distribuite care să

- dea mereu un rezultat corect.
- dea mereu același rezultat având aceleași intrări.
- funcționeze chiar dacă părți din sistem se opresc.
- funcționeze chiar dacă și părți din sistem se comportă defectuos (au bug-uri).
- funcționeze chiar dacă părți din sistem sunt compromise și devin malițioase.



# Presupunem un sistem bancar



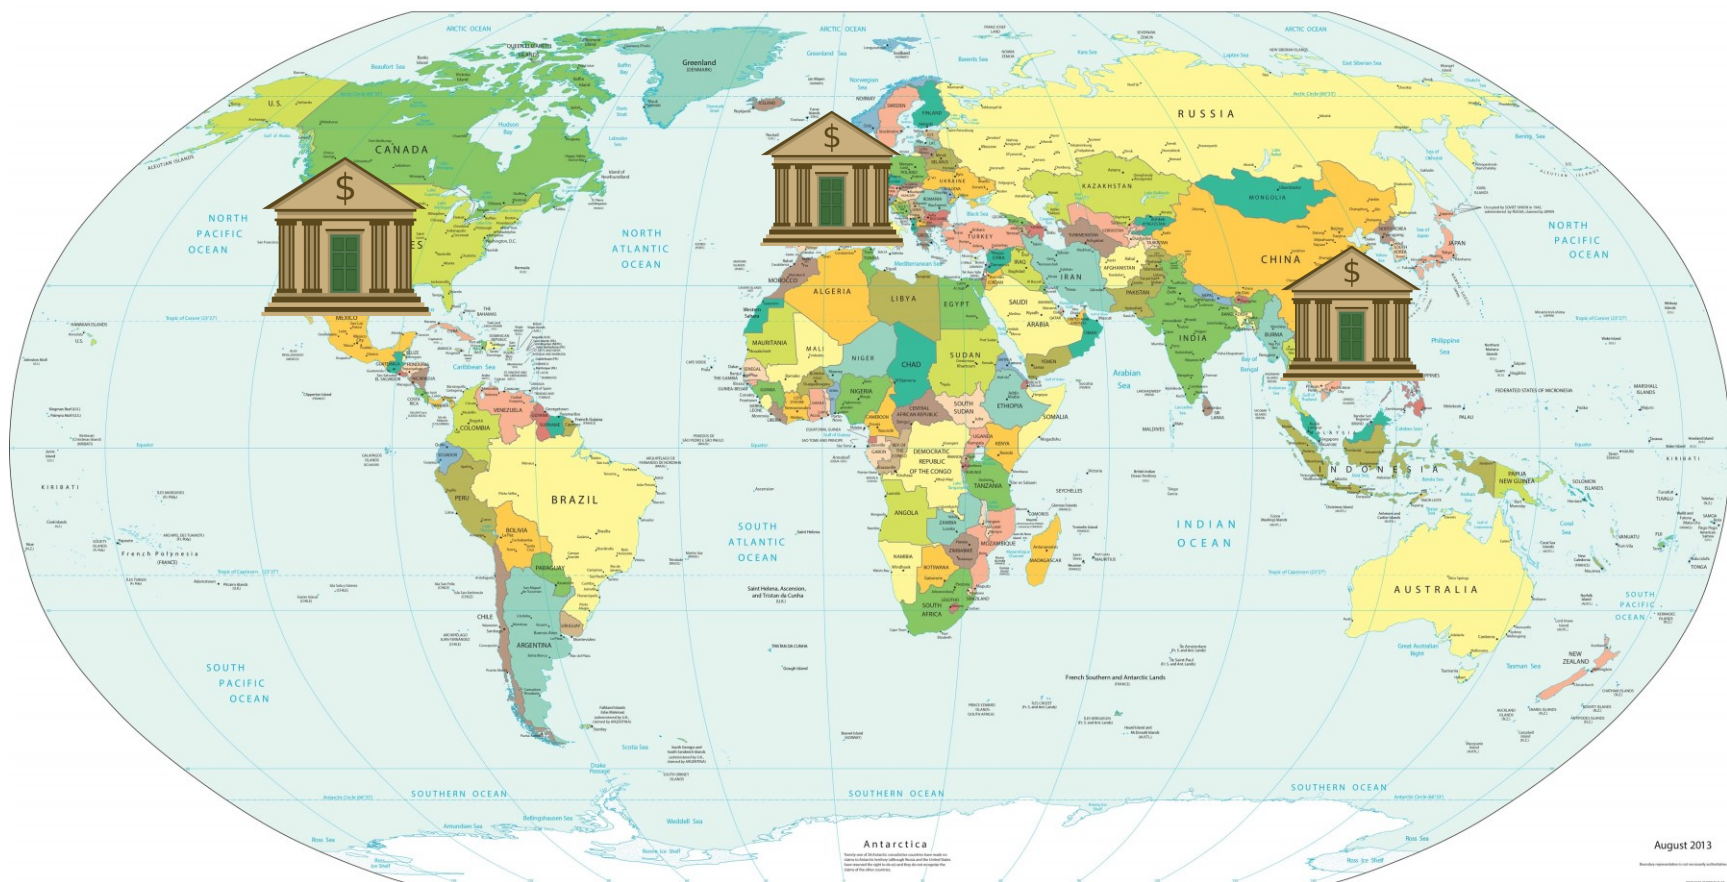


# Presupunem un sistem bancar scalabil





# Presupunem un sistem bancar scalabil distribuit global





# Presupunem un sistem bancar scalabil distribuit global

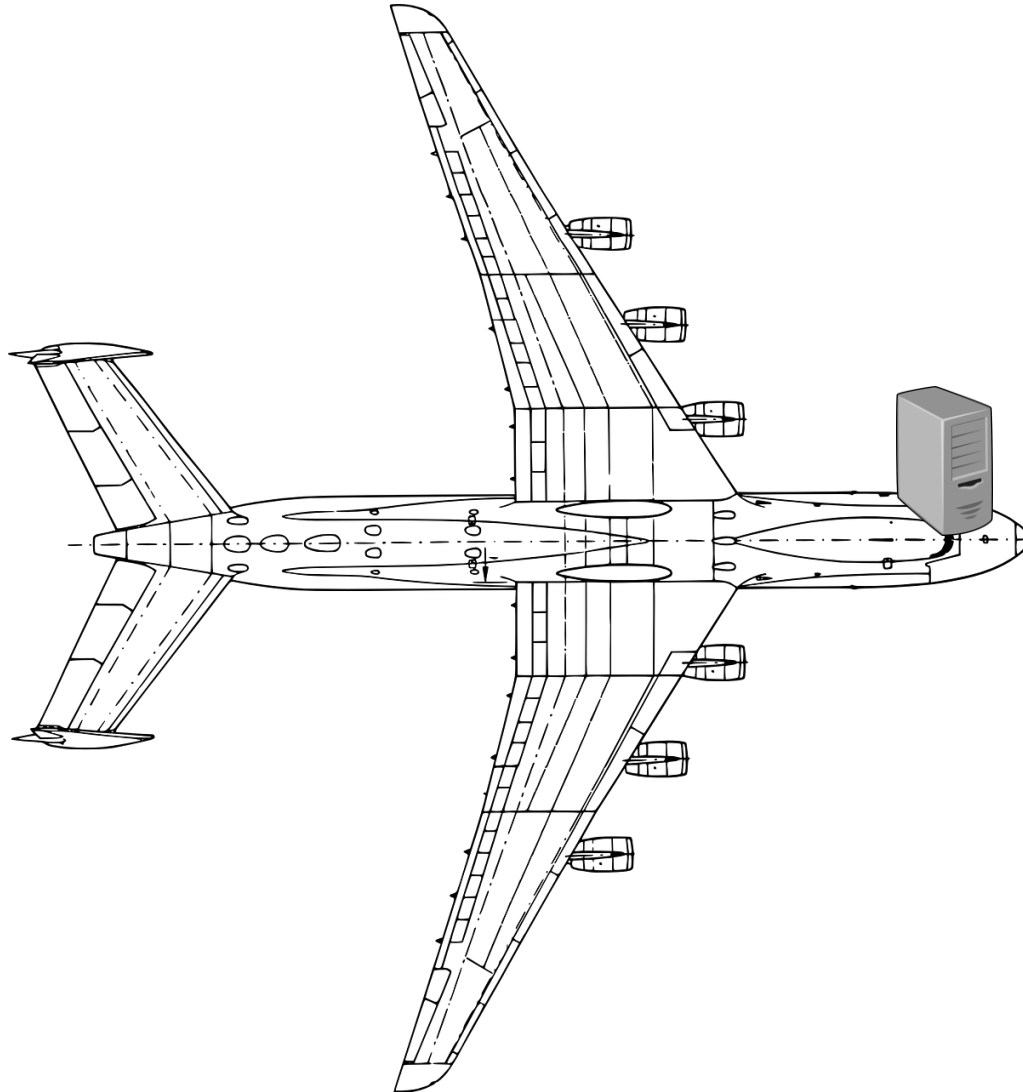


Dacă pun bani într-o bancă vreau să îi pot scoate din alta.

Dacă scot bani dintr-o bancă când alta calculează o taxă ca procent din banii mei vreau să fie calculat corect.



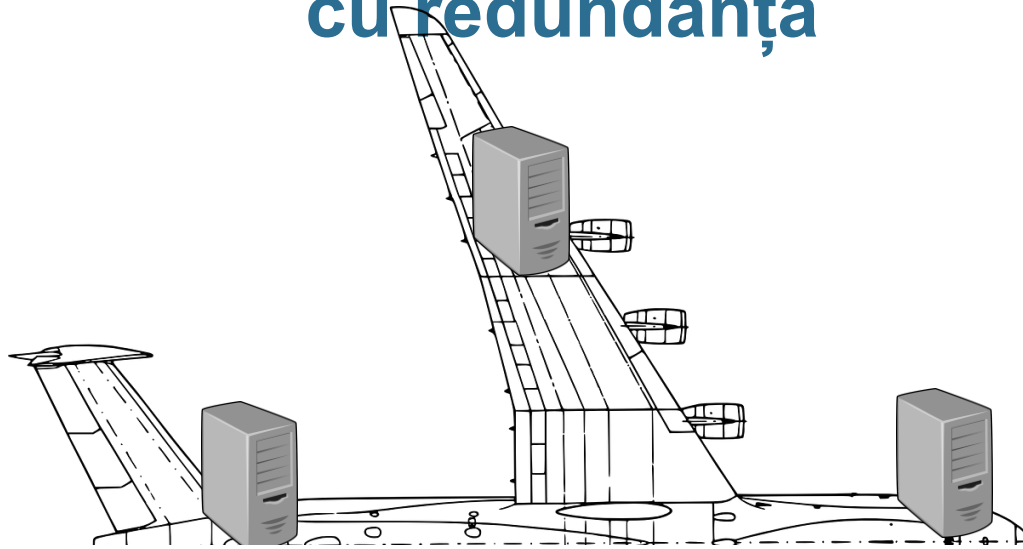
# Presupun un sistem ce controlează un avion





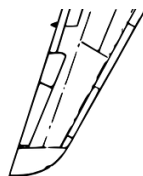


# Presupun un sistem ce controlează un avion cu redundanță



Dacă un calculator se defectează vreau ca avionul să poată funcționa.

Dacă un calculator dă comenzi greșite vreau ca celelalte să ia decizia corectă.







## Generalizare

- Am mai multe sisteme rulând un program distribuit.
- Oricărui sistem îi dau comenzi (sau le generează singur) vreau ca totalitatea lor să vadă comanda ca și executată.
- Ordinea comenzilor (chiar dacă sunt date pe sisteme diferite) este importantă.  $(A + B) / 2 \neq A / 2 + B$
- Dacă unele sisteme pică, programul distribuit trebuie să funcționeze corect în continuare.
- Dacă unele sisteme se comportă malițios, doresc ca programul distribuit să funcționeze corect în continuare.
- Toate aceste exemple sunt inexistente în programele secvențiale



## Generalizare

- Automat avem comunicare asincronă.
- Dacă comunicarea ar fi sincronă și un sistem s-ar opri, un altul ar putea și el rămâne blocat așteptând după acesta.



**Multe din probleme ar avea o rezolvare ușoară  
dacă am putea atașa fiecărei comenzi timpul  
curent**

**Dar timpul în sine e o problemă dificilă în sisteme  
distribuite**

**Dacă avem două sisteme avem două “păreri”  
despre timpul curent.**



# Cu ce protocol se sincronizează ceasul calculatoarelor?



# Sincronizarea timpilor între 2 sisteme

## Probabilistic clock synchronization

**Flaviu Cristian**

IBM Almaden Research Center, 650 Harry Road, San Jose, CA 95120, USA



**Flaviu Cristian** is a computer scientist at the IBM Almaden Research Center in San Jose, California. He received his PhD from the University of Grenoble, France, in 1979. After carrying out research in operating systems and programming methodology in France, and working on the specification, design, and verification of fault-tolerant programs in England, he joined IBM in 1982. Since then he has worked in the area of fault-tolerant distributed protocols and

systems. He has participated in the design and implementation of a highly available system prototype at the Almaden Research Center and has reviewed and consulted for several fault-tolerant distributed system designs, both in Europe and in the American divisions of IBM. He is now a technical leader in the design of a new U.S. Air Traffic Control System which must satisfy very stringent availability requirements.

**Abstract.** A probabilistic method is proposed for reading remote clocks in distributed systems subject to unbounded random communication delays.

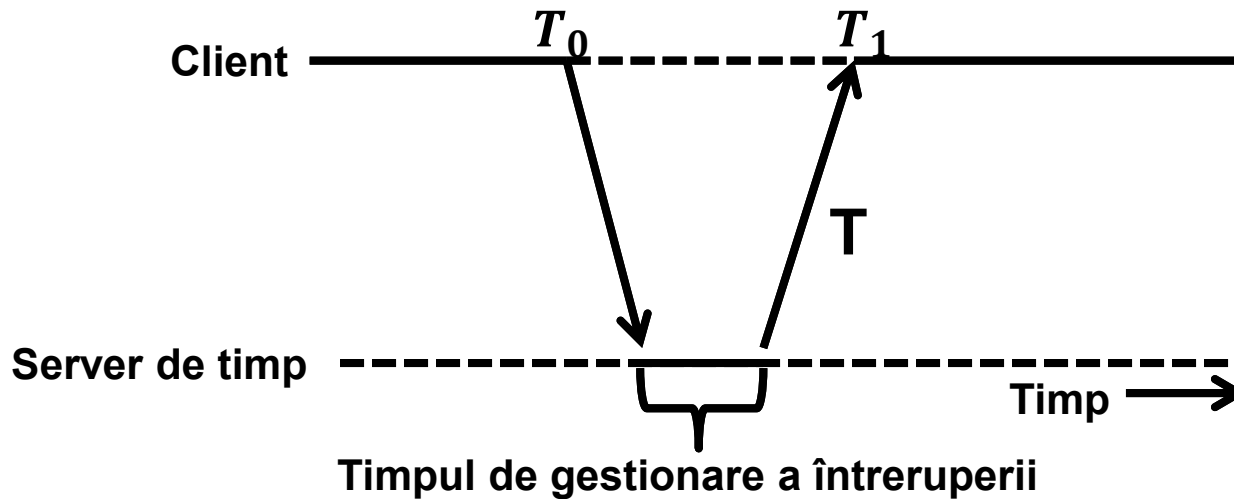
in some given maximum derivation from a time reference external to the system. *Internal* clock synchronization keeps processor clocks within some maximum relative deviation of each other. Externally synchronized clocks are also internally synchronized. The converse is not true: as time passes internally synchronized clocks can drift arbitrarily far from external time.

Clock synchronization is needed in many distributed systems. Internal clock synchronization enables one to measure the duration of distributed activities that start on one processor and terminate on another processor and to totally order distributed events in a manner that closely approximates their real time precedence. To allow exchange of information about the timing of events with other systems and users, many systems require external clock synchronization. For example external time can be used to record the occurrence of events for later analysis by humans, to instruct a system to take certain actions when certain specified (external) time deadlines occur, and to order the occurrence of related events observed by distinct sys-



# Sincronizarea timpilor între 2 sisteme

$$T + \text{RTT}/2$$





# Cu ce protocol se sincronizează ceasul calculatoarelor?

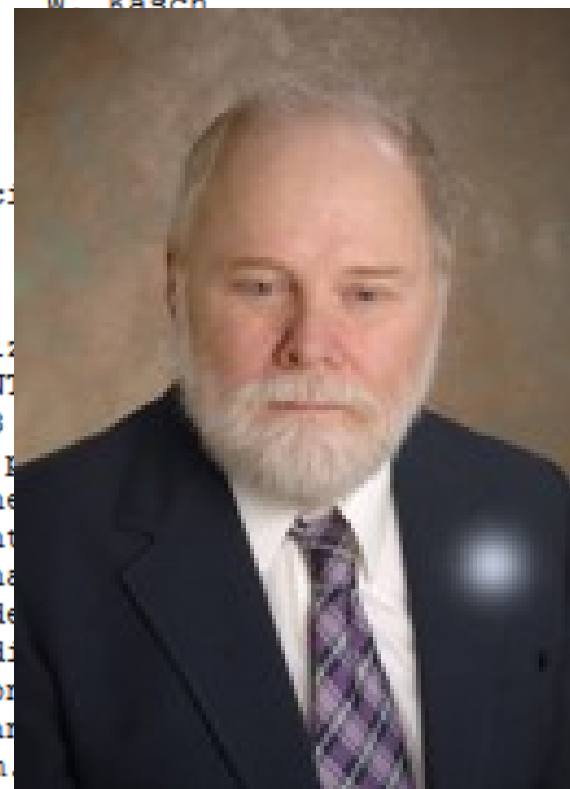
Internet Engineering Task Force (IETF)  
Request for Comments: 5905  
Obsoletes: 1305, 4330  
Category: Standards Track  
ISSN: 2070-1721

D. Mills  
U. Delaware  
J. Martin, Ed.  
ISC  
J. Burbank  
W. Kasch

Network Time Protocol Version 4: Protocol and Algorithms Specification

## Abstract

The Network Time Protocol (NTP) is widely used to synchronize computer clocks in the Internet. This document describes NTP Version 4 (NTPv4), which is backwards compatible with NTP version 3 as described in RFC 1305, as well as previous versions of the protocol. NTPv4 includes a modified protocol header to accommodate the Protocol version 6 address family. NTPv4 includes fundamental improvements in the mitigation and discipline algorithms that increase the potential accuracy to the tens of microseconds with modern workstations and fast LANs. It includes a dynamic server discovery scheme, so that in many cases, specific server configuration is not required. It corrects certain errors in the NTPv3 design and implementation and includes an optional extension mechanism.







## Ce “viteză” au calculatoarele?

- Reminder: viteză e un cuvânt extrem de nepotrivit



# În cât timp va executa următorul program?

```
int main() {  
    → int i;  
    → for (i=0; i < 1000 * 1000 * 1000; i++);  
    → return 0;  
}
```



# În cât timp va executa următorul program?

```
int main() {  
    → int i;  
    → for (i=0; i < 1000 * 1000 * 1000; i++);  
    → return 0;  
}
```

```
real    0m1.783s  
user    0m1.766s  
sys     0m0.000s
```



# Timp transmite pachet?



# Timp transmitere pachet?

```
root@Nyx:/mnt/d/Dropbox/backupServer/apd-homework/rezults# ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=120 time=15.4 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=120 time=13.0 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=120 time=15.3 ms
64 bytes from 8.8.8.8: icmp_seq=4 ttl=120 time=14.5 ms
64 bytes from 8.8.8.8: icmp_seq=5 ttl=120 time=32.5 ms
64 bytes from 8.8.8.8: icmp_seq=6 ttl=120 time=15.2 ms
```



# Scala

Termen	Not	Multiplu	$10^x$	
Tera	T	1 000 000 000 000	12	
Giga	G	1 000 000 000	9	
Mega	M	1 000 000	6	
kilo	k	1 000	3	
UNITATE		1	0	
mili	m	0. 001	-3	Trimitere pachet
micro	$\mu$	0. 000 001	-6	
nano	n	0. 000 000 001	-9	Executare operație
pico	p	0. 000 000 000 001	-12	



**Pentru a avea încrede în obținerea proprietăților  
dorite trebuie să validăm matematic sistemele  
noastre**

**Pentru a valida un program matematic acesta  
trebuie modelat**





# Modelarea unui sistem simplu/secvențial?



# Modelarea unui sistem simplu/secvențial?

First Draft of a Report  
on the EDVAC

by

John von Neumann





# Modelarea unui sistem simplu/secvențial?

Un set de stări.

Operațiunile sau acțiunile modifică starea



# Modelarea unui sistem distribuit?



# Modelarea unui sistem distribuit?

**Obligatoriu trebuie să ia în calcul comunicarea**

**Se pot executa milioane de operații până la primirea unui mesaj de la un alt sistem**



# Modelarea unui sistem distribuit: CSP

Programming  
Techniques

S. L. Graham, R. L. Rivest  
Editors

## Communicating Sequential Processes

C.A.R. Hoare  
The Queen's University  
Belfast, Northern Ireland

**This paper suggests that input and output are basic primitives of programming and that parallel composition of communicating sequential processes is a fundamental program structuring method. When combined with a development of Dijkstra's guarded command, these concepts are surprisingly versatile. Their use is illustrated by sample solutions of a variety of familiar programming exercises.**

**Key Words and Phrases:** programming, programming languages, programming primitives, program structures, parallel programming, concurrency, input, output, guarded commands, nondeterminacy, coroutines, procedures, multiple entries, multiple exits, classes, data representations, recursion, conditional critical regions, monitors, iterative arrays

**CR Categories:** 4.20, 4.22, 4.32

grams, three basic constructs have received widespread recognition and use: A repetitive construct (e.g. the **while** loop), an alternative construct (e.g. the conditional **if..then..else**), and normal sequential program composition (often denoted by a semicolon). Less agreement has been reached about the design of other important program structures, and many suggestions have been made: Subroutines (Fortran), procedures (Algol 60 [15]), entries (PL/I), coroutines (UNIX [17]), classes (SIMULA 67 [5]), processes and monitors (CPL [12]), forms (ALPHA [13]),

The traditional stored-program computer has been designed primarily for single sequential programs. Program speed has led to the introduction of parallelism. An attempt has been made to design a computer for a programmer, either by having function units of the computer (in an I/O control package generating system). However, technology suggests that a computer should be structured from a number of processors (each with its own powerful, capacious, reliable machine which is disguised).

In order to use such a computer for a task, the component processes must communicate and to synchronize methods of achieving this. The adopted method of communication is the updating of a common state and many machine code severe problems in the com-





# Modelarea unui sistem distribuit: Rețele Petri

K o m m u n i k a t i o n .  
m i t  
A u t o m a t e n

Von der Fakultät für Mathematik und Physik  
der Technischen Hochschule Darmstadt

zur Erlangung des Grades eines  
Doktors der Naturwissenschaften  
(Dr. rer.nat.)

genehmigte  
Dissertation

vorgelegt von  
C a r l A d a m P e t r i







# Modelarea unui sistem distribuit: Mașini de stări

Operating  
Systems

R. Stockton Gaines  
Editor

## Time, Clocks, and the Ordering of Events in a Distributed System

Leslie Lamport  
Massachusetts Computer Associates, Inc.

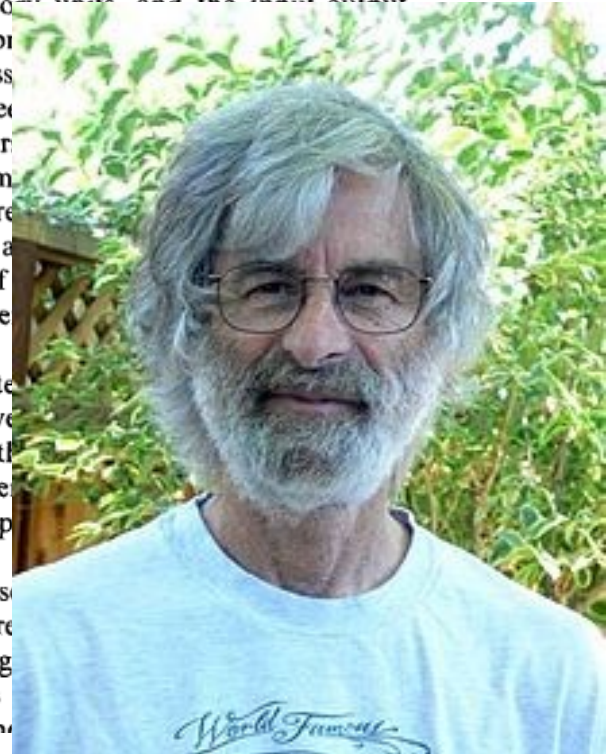
The concept of one event happening before another in a distributed system is examined, and is shown to define a partial ordering of the events. A distributed algorithm is given for synchronizing a system of logical clocks which can be used to totally order the events. The use of the total ordering is illustrated with a method for solving synchronization problems. The algorithm is then specialized for synchronizing physical clocks, and a bound is derived on how far out of synchrony the clocks can become.

A distributed system consists of a collection of distinct processes which are spatially separated, and which communicate with one another by exchanging messages. A network of interconnected computers, such as the ARPA net, is a distributed system. A single computer can also be viewed as a distributed system in which the central control unit, the memory, and the input/output channels are separate processes. The time between two events is compared to the time between two events.

We will concern ourselves with the problem of ordering spatially separated events. The remarks will apply more generally to a distributed system on a multiprocessor system on a multiprocessor system. The problems are similar to those of the unpredictable order of events.

In a distributed system, we say that one of two events "happened before" is the case if the first event is of the events in the system. This often arises because people often arise because people and its implications.

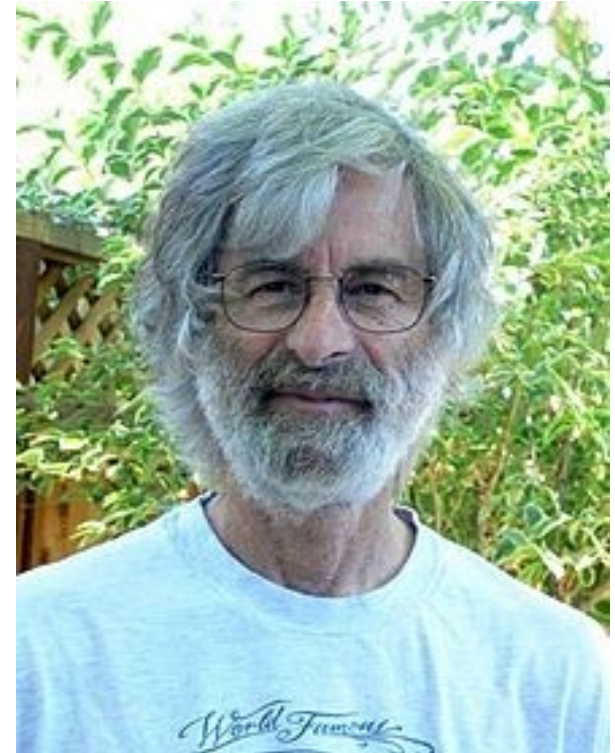
In this paper, we discuss the problem of ordering by the "happened before" algorithm for extending of all the events. This mechanism for implementing





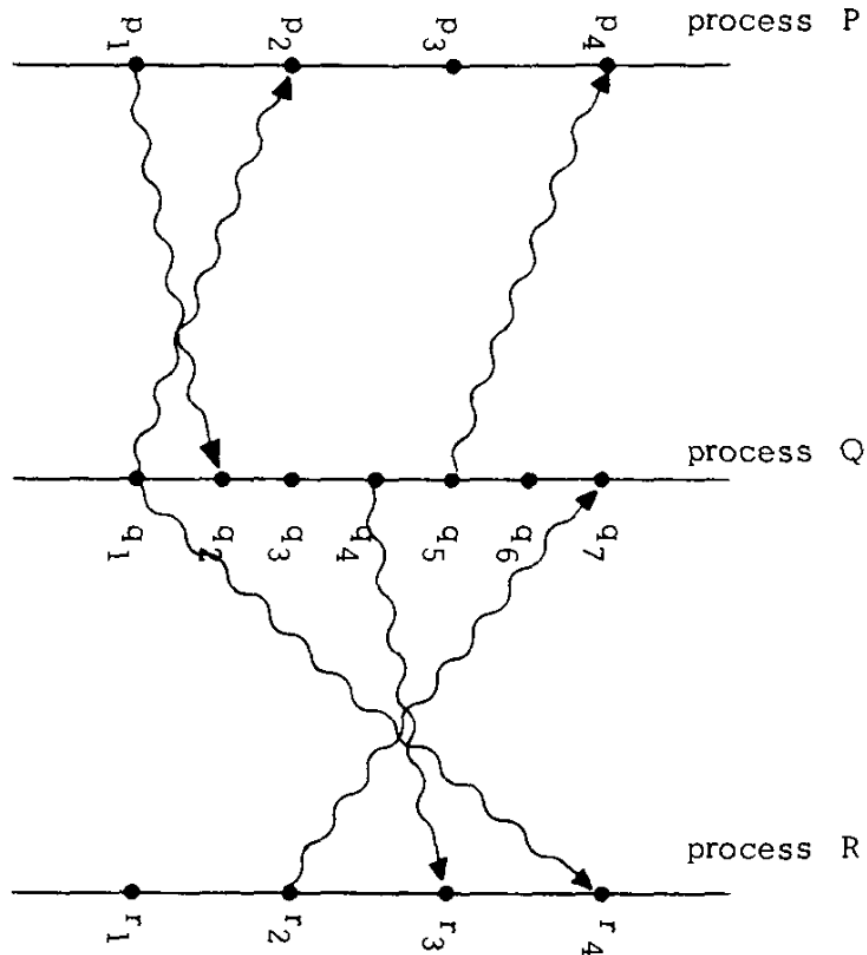
# Leslie Lamport

- Majoritatea problemelor/soluțiilor prezentate în acest curs sunt construite de Leslie Lamport.
- Câștigător premiul **Turing 2013**
- Are una din cele mai citate lucrări
- Ne este contemporan și are multiple înregistrări pe Youtube scurte:  
[1](#) [2](#) [3](#) [4](#) [5](#)
- Și un [interviu lung](#) în [două părți](#)





# Lamport – ceasuri logice, ordonare parțială





## Lamport – ceasuri logice, ordonare parțială

IR1. Each process  $P_i$  increments  $C_i$  between any two successive events.

IR2. (a) If event  $a$  is the sending of a message  $m$  by process  $P_i$ , then the message  $m$  contains a timestamp  $T_m = C_i\langle a \rangle$ . (b) Upon receiving a message  $m$ , process  $P_j$  sets  $C_j$  greater than or equal to its present value and greater than  $T_m$ .

*Clock Condition.* For any events  $a, b$ :

if  $a \rightarrow b$  then  $C\langle a \rangle < C\langle b \rangle$ .

**Invers NU este adevărat**



## Lamport – ordonare totală $\Rightarrow$

- Putem să definim o relație de ordine totală folosindu-ne de id-ul procesului.
- Dacă două evenimente au același C atunci le comparăm și după id-ul procesului.



## Lamport – mutex distribuit

1. To request the resource, process  $P_i$  sends the message  $T_m:P_i$  *requests resource* to every other process, and puts that message on its request queue, where  $T_m$  is the timestamp of the message.
2. When process  $P_j$  receives the message  $T_m:P_i$  *requests resource*, it places it on its request queue and sends a (timestamped) acknowledgment message to  $P_i$ .<sup>5</sup>
3. To release the resource, process  $P_i$  removes any  $T_m:P_i$  *requests resource* message from its request queue and sends a (timestamped)  $P_i$  *releases resource* message to every other process.





## Lamport – mutex distribuit

4. When process  $P_j$  receives a  $P_i$  *releases resource* message, it removes any  $T_m:P_i$  *requests resource* message from its request queue.

5. Process  $P_i$  is granted the resource when the following two conditions are satisfied: (i) There is a  $T_m:P_i$  *requests resource* message in its request queue which is ordered before any other request in its queue by the relation  $\Rightarrow$ . (To define the relation “ $\Rightarrow$ ” for messages, we identify a message with the event of sending it.) (ii)  $P_i$  has received a message from every other process time-stamped later than  $T_m$ .<sup>6</sup>







# Vector Clocks

## Timestamps in Message-Passing Systems That Preserve the Partial Ordering

*Colin J. Fidge*

*Department of Computer Science, Australian National University, Canberra, ACT.*

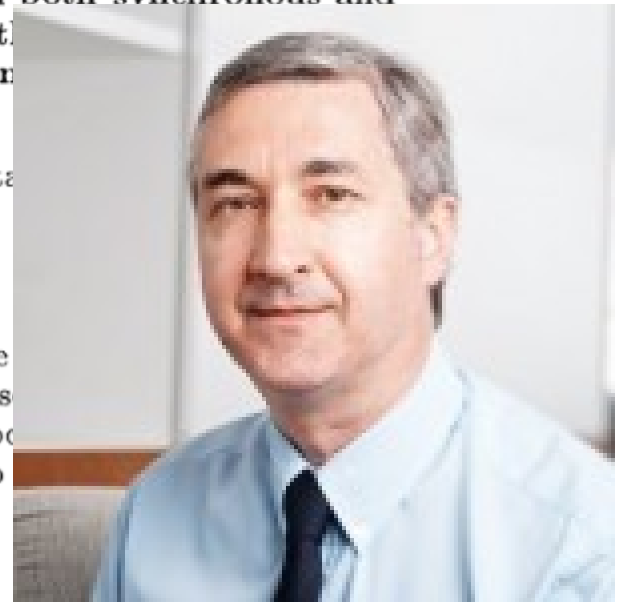
### ABSTRACT

Timestamping is a common method of totally ordering events in concurrent programs. However, for applications requiring access to the global state, a total ordering is inappropriate. This paper presents algorithms for timestamping events in both synchronous and asynchronous message-passing programs that allow for access to the global state without being herent in a parallel system. The algorithms do not change the communication semantics and require a central timestamp issuing authority.

**Keywords and phrases:** concurrent programming, message-passing, timestamps  
**CR categories:** D.1.3

### INTRODUCTION

A fundamental problem in concurrent programming is determining the order of events that occurred. An obvious solution is to attach a number representing the order of events to each event. This assumes that each process maintains a permanent record of the execution of each event. This assumes that each process has a global clock, but practical parallel systems, by their very nature, make it difficult to





# Vector Clocks

## Virtual Time and Global States of Distributed Systems \*

Friedemann Mattern <sup>†</sup>

Department of Computer Science, University of Kaiserslautern  
D 6750 Kaiserslautern, Germany

### Abstract

*A distributed system can be characterized by the fact that the global state is distributed and that a common time base does not exist. However, the notion of time is an important concept in every day life of our decentralized “real world” and helps to solve problems like getting a consistent population census or determining the potential causality between events. We argue that a linearly ordered structure of time is not (always) adequate for distributed systems and propose a generalized non-standard model of time which consists of vectors*

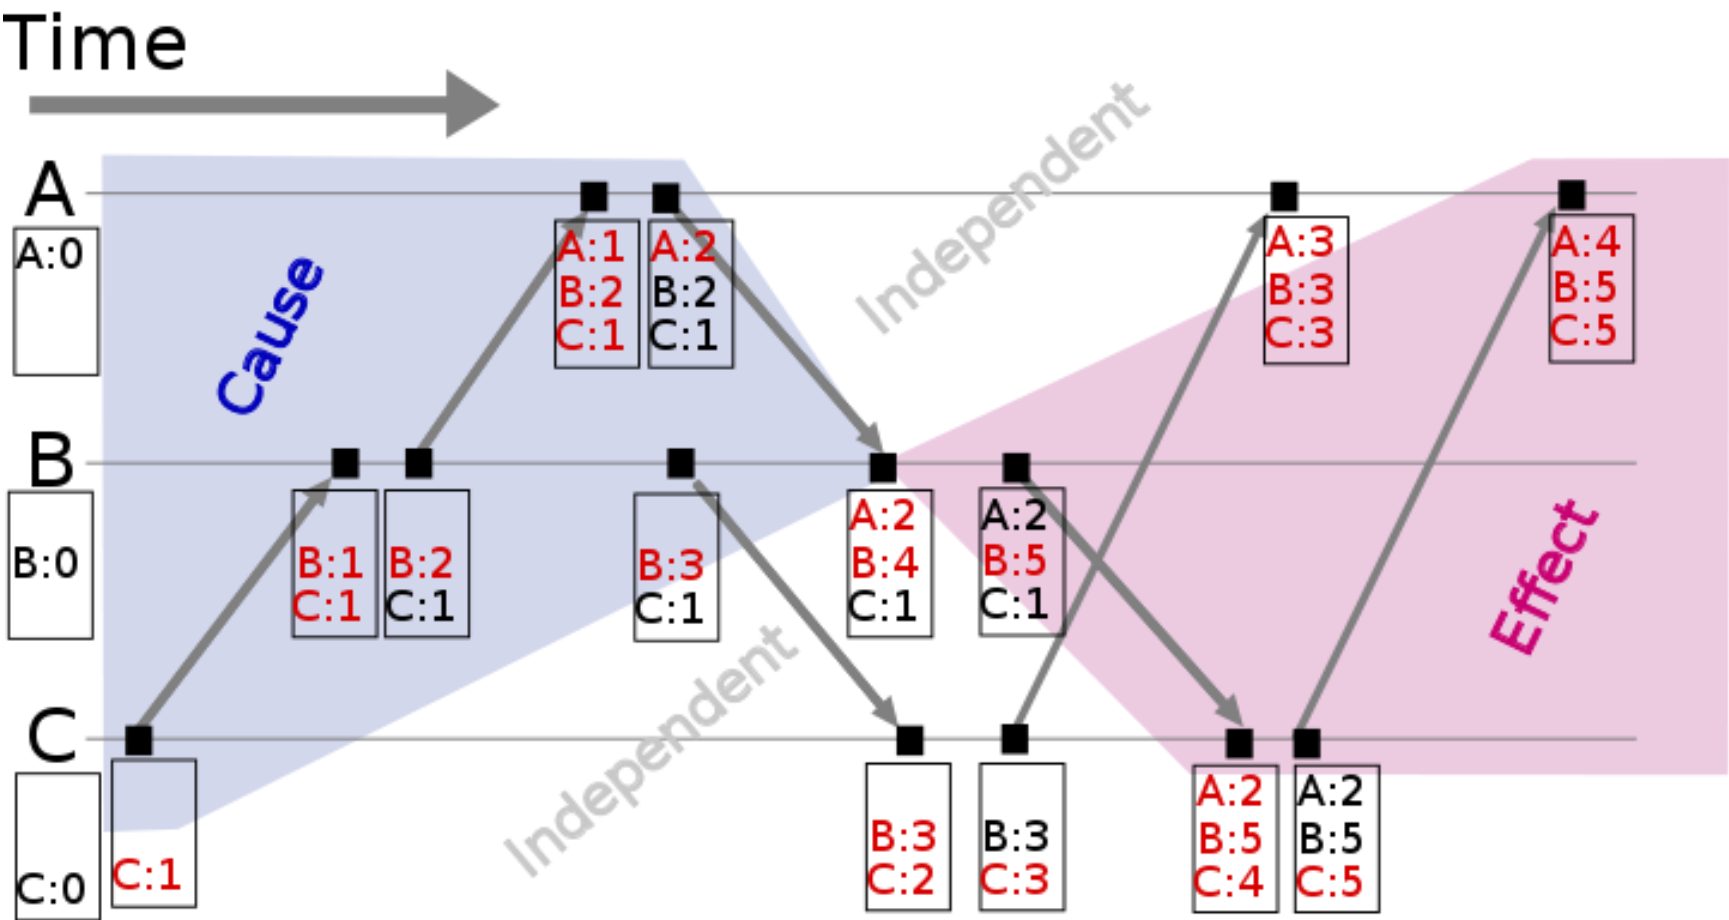
view of an idealized access to all processes.

The fact that a view of the global state does not exist is the cause of many problems in distributed systems. Consensus, replication, and database systems, deadlock detection, and concurrency control are difficult to solve in a distributed environment. In a physical centralized environment, these problems are found to be wrong. In a distributed environment, these problems are found to be wrong.





# Vector Clocks







# Problema Generalilor Bizantini

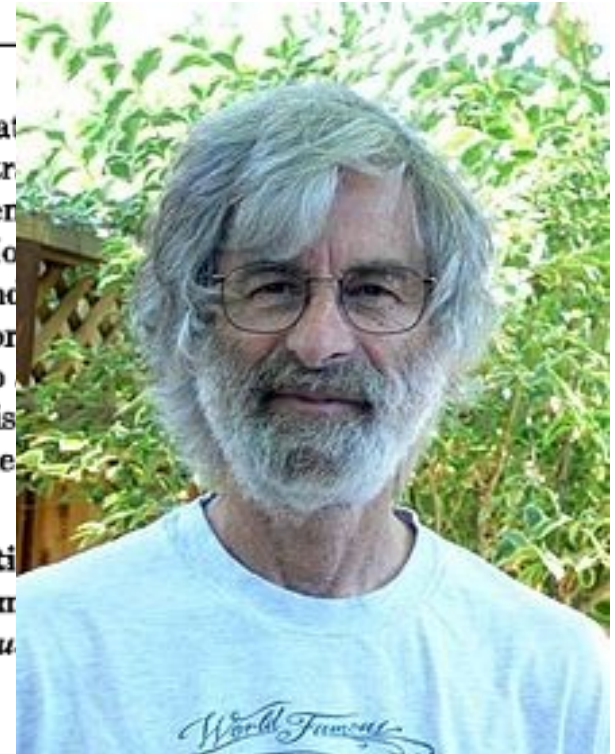
## The Byzantine Generals Problem

LESLIE LAMPORT, ROBERT SHOSTAK, and MARSHALL PEASE  
SRI International

Reliable computer systems must handle malfunctioning components that may fail arbitrarily and return inconsistent values. This situation can be expressed abstractly as the Byzantine Generals Problem. In this problem, a set of  $n$  generals surround an enemy city. They must agree upon a common battle plan. However, some of the generals may be traitors who will try to confuse the others. The problem is to find an algorithm that allows the loyal generals to reach agreement. It is shown that, using only oral messages, the problem is solvable if and only if more than two-thirds of the generals are loyal; so, for example, with three generals, two must be loyal. With unforgeable written messages, the problem is solvable for any number of loyal generals and possible traitors. Applications of the solutions to reliable computer systems are discussed.

Categories and Subject Descriptors: C.2.4. [Computer-Communications Systems]—*network operating systems*; D.4.4 [Operating Systems]: Communications—*network communication*; D.4.5 [Operating Systems]: Reliability—*fault tolerance*

General Terms: Algorithms, Reliability





# Soluție Generali Bizantini – fără semnături

- Există doar dacă mai mult de două treimi din generali sunt onești.
  - **NU** există soluție pentru 3 sau mai puțini generali.
- Necesită multă comunicație, mesajele de la un general sunt transmise la toți ceilalți într-o formă recursivă.





# Soluție Generali Bizantini – cu semnătură

- (a) A loyal general's signature cannot be forged, and any alteration of the contents of his signed messages can be detected.
- (b) Anyone can verify the authenticity of a general's signature.

Funcționează dacă sunt minim doi generali corecți.  
Înseamnă că există soluție pentru 3 generali.

Our algorithm assumes a function *choice* which is applied to a set of orders to obtain a single one. The only requirements we make for this function are

1. If the set  $V$  consists of the single element  $v$ , then  $choice(V) = v$ .
2.  $choice(\emptyset) = \text{RETREAT}$ , where  $\emptyset$  is the empty set.





# Soluție Generali Bizantini – cu semnătură

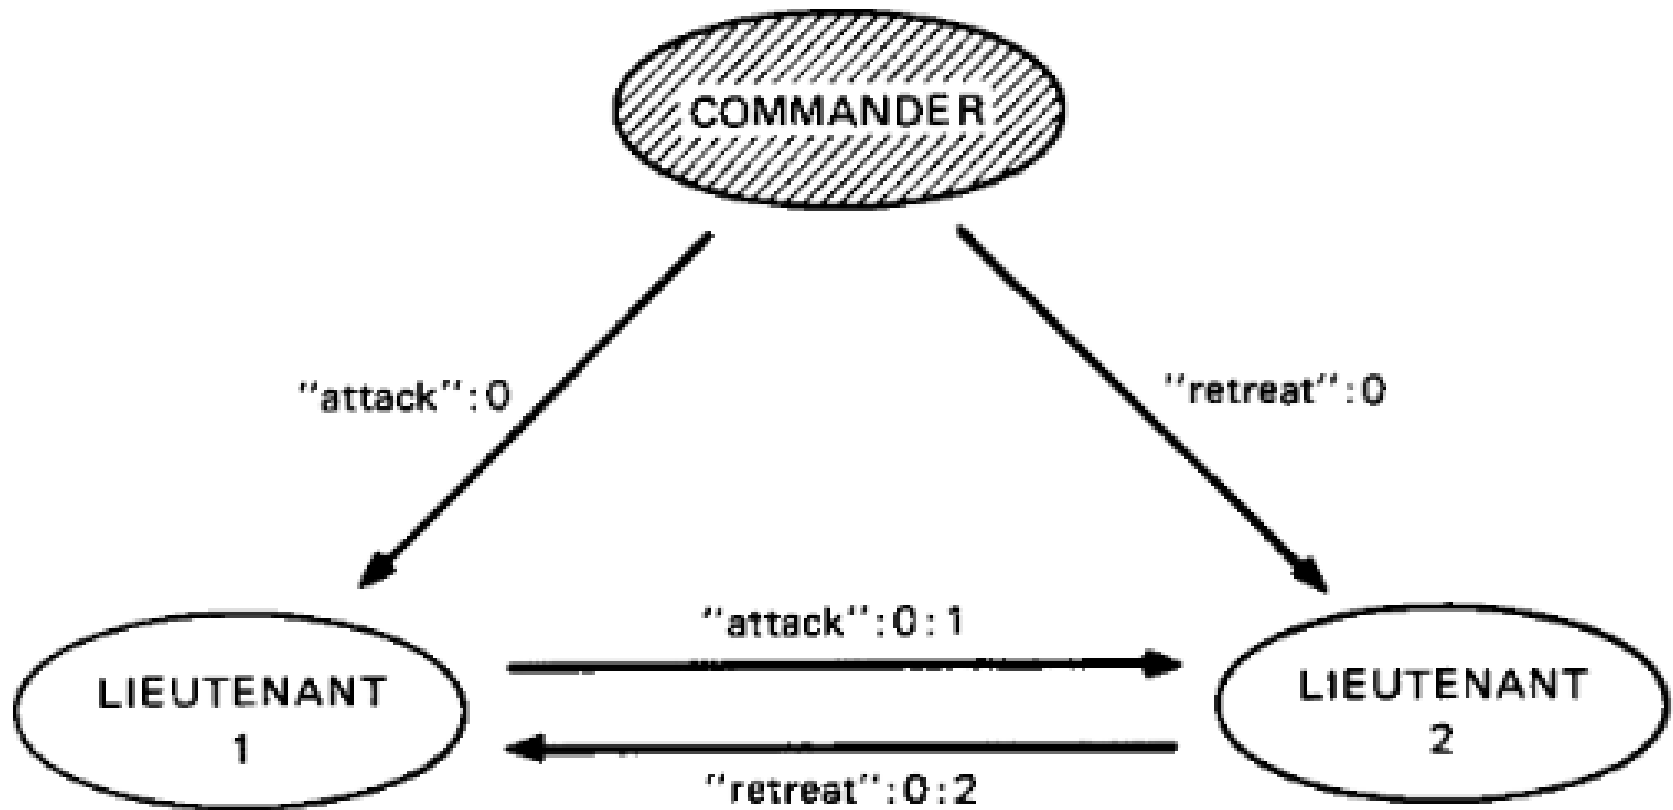
*Algorithm SM(m).*

Initially  $V_i = \emptyset$ .

- (1) The commander signs and sends his value to every lieutenant.
- (2) For each  $i$ :
  - (A) If Lieutenant  $i$  receives a message of the form  $v:0$  from the commander and he has not yet received any order, then
    - (i) he lets  $V_i$  equal  $\{v\}$ ;
    - (ii) he sends the message  $v:0:i$  to every other lieutenant.
  - (B) If Lieutenant  $i$  receives a message of the form  $v:0:j_1:\dots:j_k$  and  $v$  is not in the set  $V_i$ , then
    - (i) he adds  $v$  to  $V_i$ ;
    - (ii) if  $k < m$ , then he sends the message  $v:0:j_1:\dots:j_k:i$  to every lieutenant other than  $j_1, \dots, j_k$ .
- (3) For each  $i$ : When Lieutenant  $i$  will receive no more messages, he obeys the order *choice*( $V_i$ ).



# Soluție Generali Bizantini – cu semnătură







# Consens rezistent la Faults - Paxos

## The Part-Time Parliament

LESLIE LAMPORT

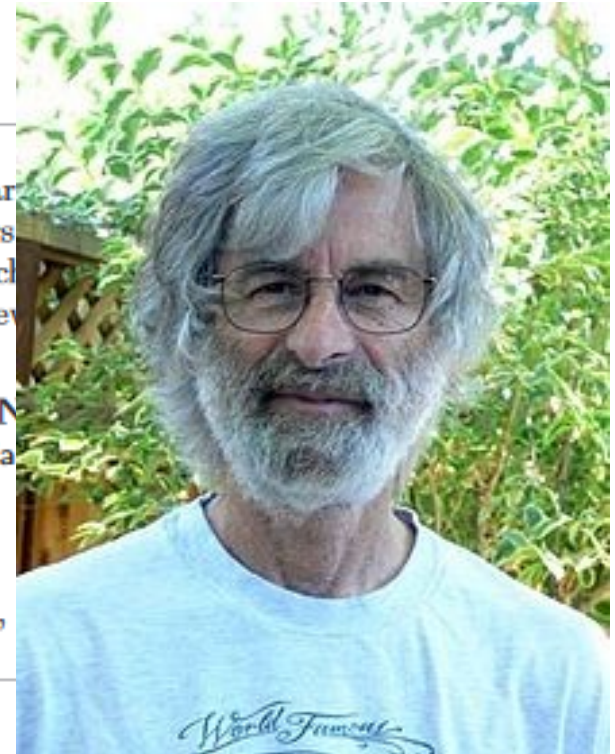
Digital Equipment Corporation

Recent archaeological discoveries on the island of Paxos reveal that the part-time legislators, despite the peripatetic propensity of its part-time legislators. The legislators keep copies of the parliamentary record, despite their frequent forays from the cliffs and the fulness of their messengers. The Paxos parliament's protocol provides a new twist on the state-machine approach to the design of distributed systems.

Categories and Subject Descriptors: C2.4 [Computer-Communications Networks]—*Network operating systems*; D4.5 [Operating Systems]: Reliability; J.1 [Administrative Data Processing]: Government

General Terms: Design, Reliability

Additional Key Words and Phrases: State machines, three-phase commit,





# Paxos

- Phase 1.** (a) A proposer selects a proposal number  $n$  and sends a *prepare* request with number  $n$  to a majority of acceptors.
- (b) If an acceptor receives a *prepare* request with number  $n$  greater than that of any *prepare* request to which it has already responded, then it responds to the request with a promise not to accept any more proposals numbered less than  $n$  and with the highest-numbered proposal (if any) that it has accepted.



# Paxos

- Phase 2.** (a) If the proposer receives a response to its *prepare* requests (numbered  $n$ ) from a majority of acceptors, then it sends an *accept* request to each of those acceptors for a proposal numbered  $n$  with a value  $v$ , where  $v$  is the value of the highest-numbered proposal among the responses, or is any value if the responses reported no proposals.
- (b) If an acceptor receives an *accept* request for a proposal numbered  $n$ , it accepts the proposal unless it has already responded to a *prepare* request having a number greater than  $n$ .



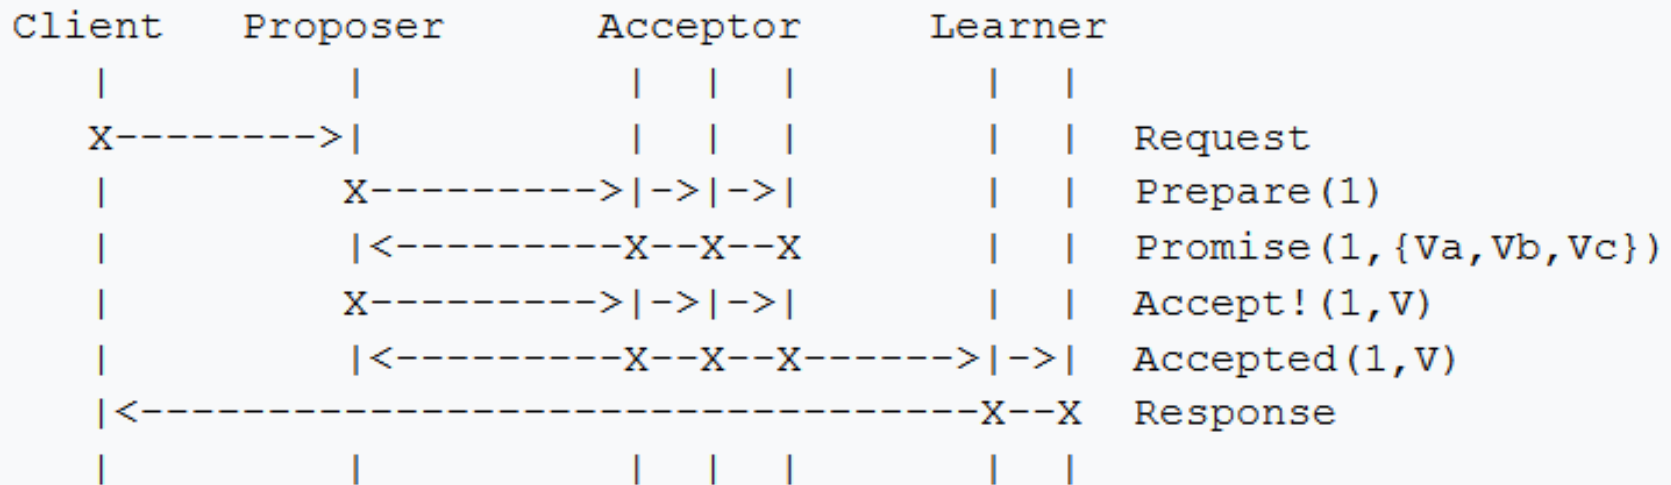
# Paxos

## 2.3 Learning a Chosen Value

To learn that a value has been chosen, a learner must find out that a proposal has been accepted by a majority of acceptors. The obvious algorithm is to have each acceptor, whenever it accepts a proposal, respond to all learners, sending them the proposal. This allows learners to find out about a chosen value as soon as possible, but it requires each acceptor to respond to each learner—a number of responses equal to the product of the number of acceptors and the number of learners.



# Paxos







# Paxos

Client	Proposer	Acceptor	Learner	
X----->				Request
	X-----> -> ->			Prepare(1)
		!		!! FAIL !!
	<-----X--X			Promise(1, {Va, Vb, null})
	X-----> ->			Accept!(1, V)
	<-----X--X-----> ->			Accepted(1, V)
<-----X--X				Response