



Sisteme Tolerante la Defecte Consens în Sisteme Distribuite 2

Lect. Dr. Ing. Cristian Chilipirea – cristian.chilipirea@mta.ro







Consensus resistant to Faults - Paxos

The Part-Time Parliament

LESLIE LAMPORT


Digital Equipment Corporation

Recent archaeological discoveries on the island of Paxos reveal that the people of Paxos, despite the peripatetic propensity of its part-time legislators. The legislators have preserved copies of the parliamentary record, despite their frequent forays from the island in the fullness of their messengers. The Paxos parliament's protocol provides a new paradigm for the state-machine approach to the design of distributed systems.

Categories and Subject Descriptors: C2.4 [Computer-Communications Systems]—*Network operating systems*; D4.5 [Operating Systems]: Reliability; J.1 [Administrative Data Processing]: Government

General Terms: Design, Reliability

Additional Key Words and Phrases: State machines, three-phase commit





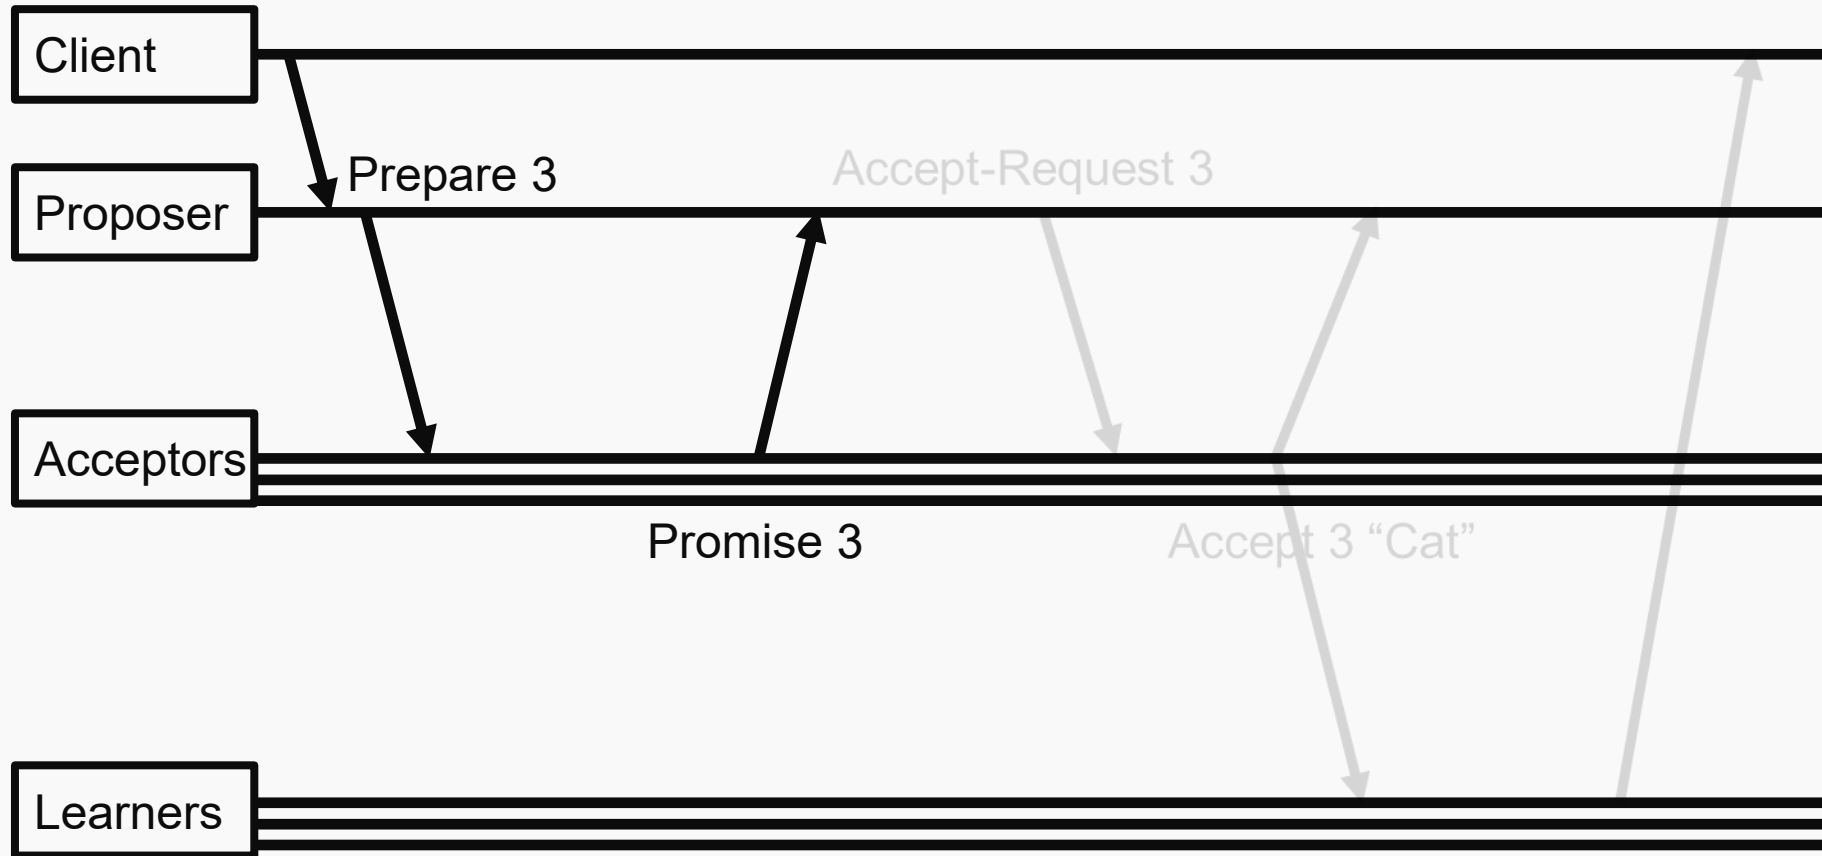
Paxos

Phase 1. (a) A proposer selects a proposal number n and sends a *prepare* request with number n to a majority of acceptors.

(b) If an acceptor receives a *prepare* request with number n greater than that of any *prepare* request to which it has already responded, then it responds to the request with a promise not to accept any more proposals numbered less than n and with the highest-numbered proposal (if any) that it has accepted.



Paxos





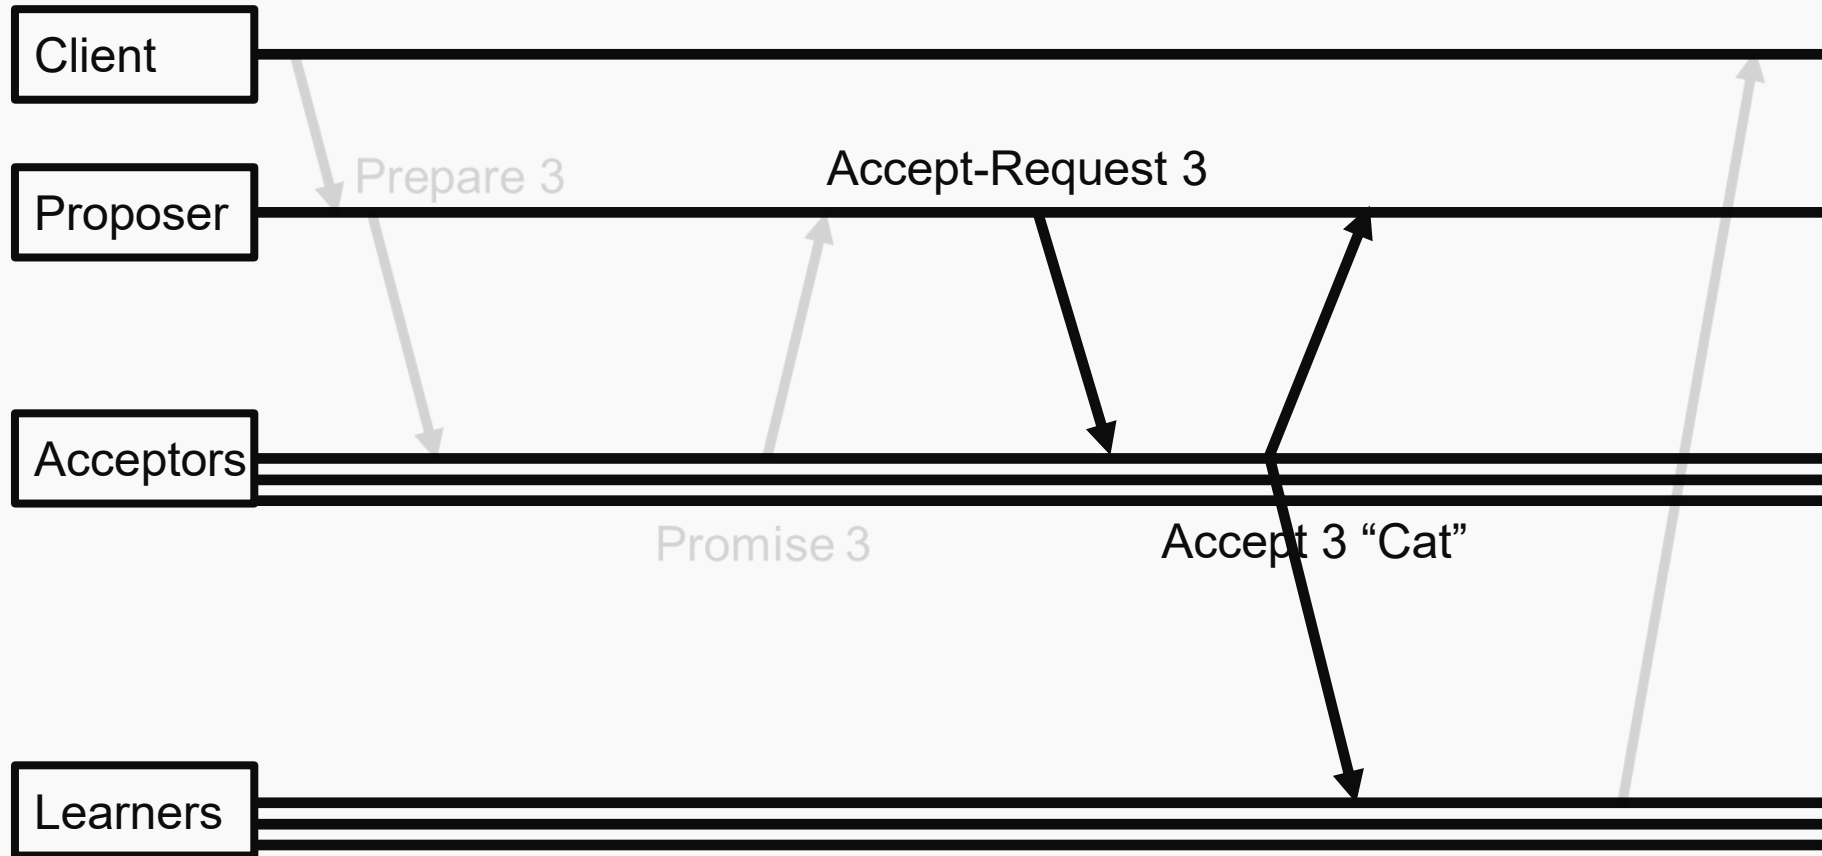
Paxos

Phase 2. (a) If the proposer receives a response to its *prepare* requests (numbered n) from a majority of acceptors, then it sends an *accept* request to each of those acceptors for a proposal numbered n with a value v , where v is the value of the highest-numbered proposal among the responses, or is any value if the responses reported no proposals.

(b) If an acceptor receives an *accept* request for a proposal numbered n , it accepts the proposal unless it has already responded to a *prepare* request having a number greater than n .



Paxos





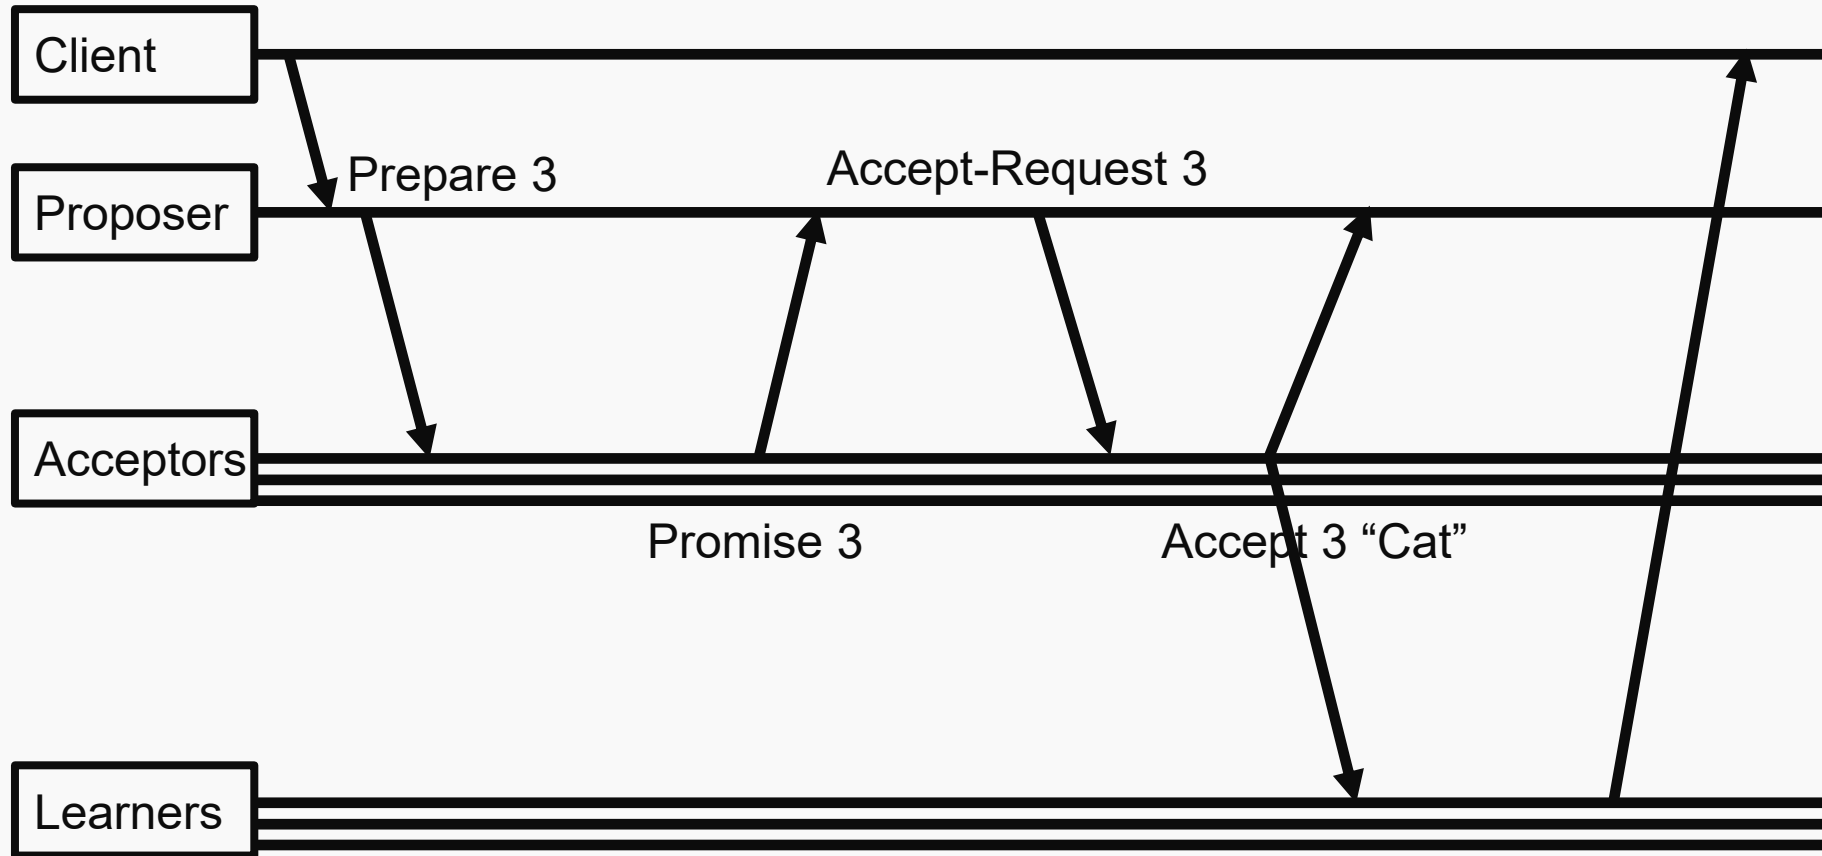
Paxos

2.3 Learning a Chosen Value

To learn that a value has been chosen, a learner must find out that a proposal has been accepted by a majority of acceptors. The obvious algorithm is to have each acceptor, whenever it accepts a proposal, respond to all learners, sending them the proposal. This allows learners to find out about a chosen value as soon as possible, but it requires each acceptor to respond to each learner—a number of responses equal to the product of the number of acceptors and the number of learners.

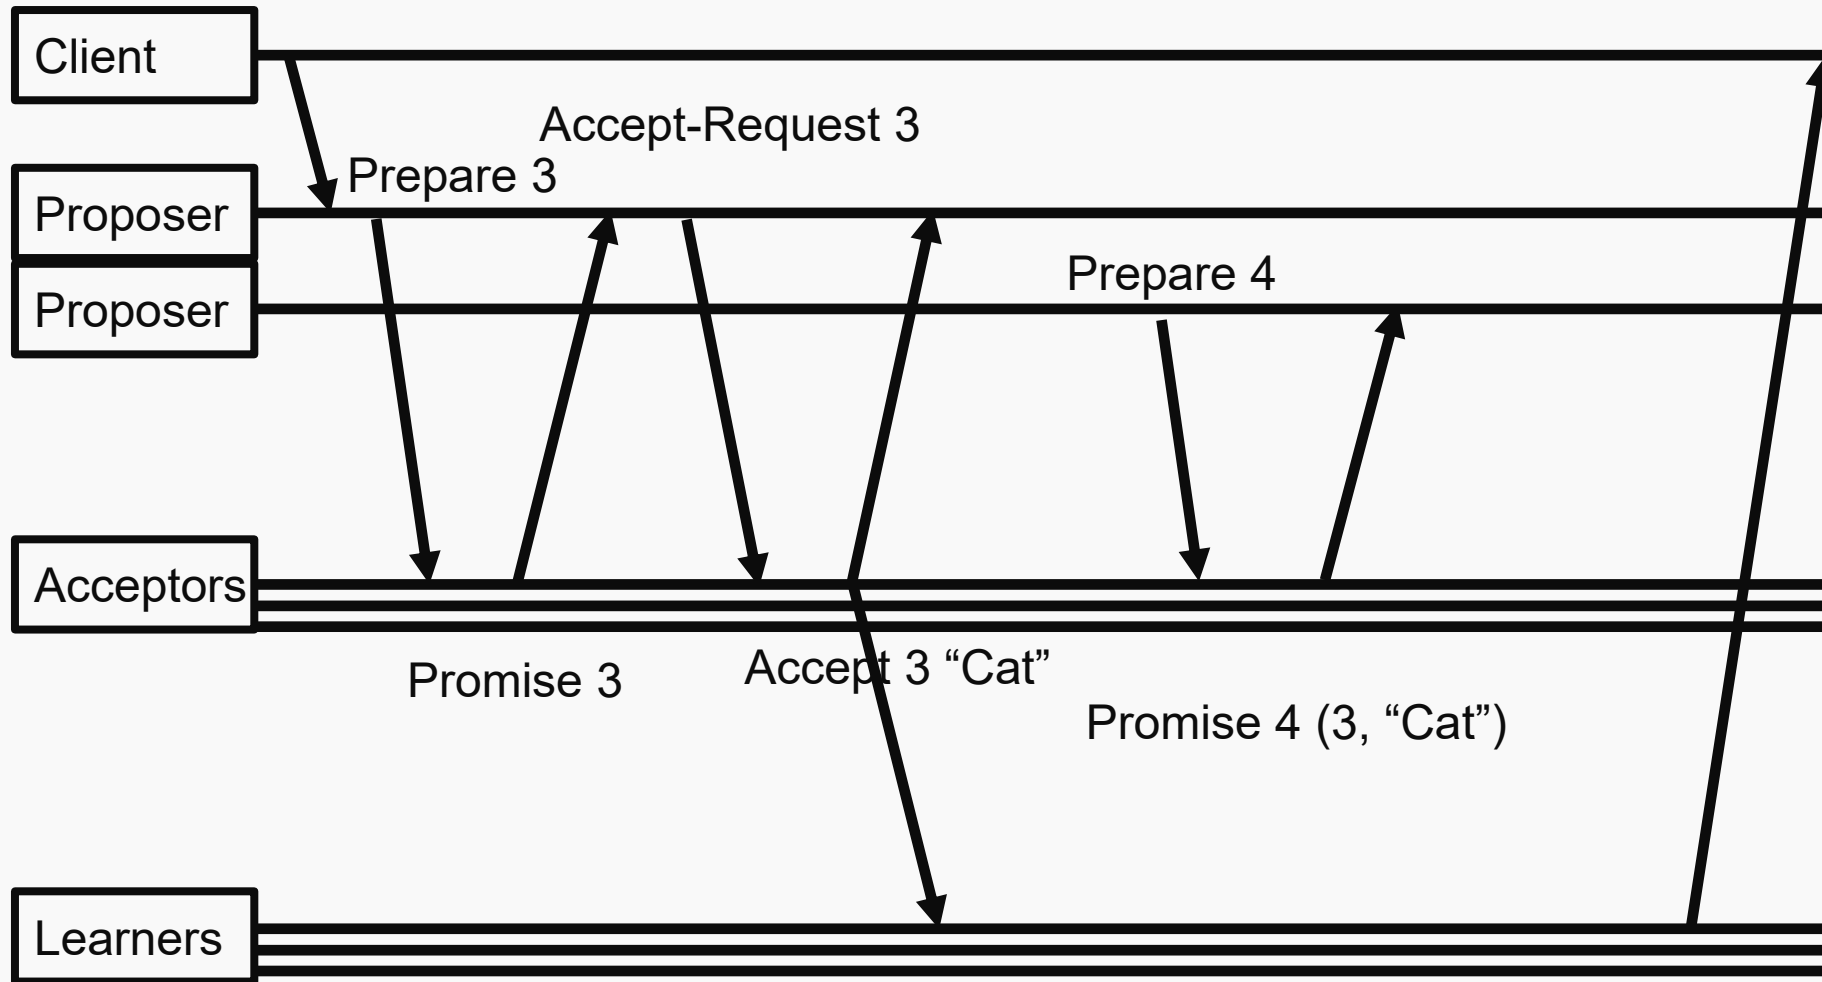


Paxos





Paxos







CAP Theorem

Brewer's Conjecture and the Feasibility of Consistent, Available, Partition-Tolerant Web Services

Seth Gilbert and Nancy Lynch
Laboratory for Computer Science
Massachusetts Institute of Technology
Cambridge, MA 02139
sethg@mit.edu, lynch@theory.lcs.mit.edu

Abstract

When designing distributed web services, there are three properties desired: consistency, availability, and partition tolerance. It is impossible to have all three. In this note, we prove this conjecture in the asynchronous model, then discuss solutions to this dilemma in the partially synchronous model.

1 Introduction

At PODC 2000, Brewer¹, in an invited talk [2], made the following conjecture: it is impossible for a web service to provide the following three guarantees:

- Consistency
- Availability
- Partition-tolerance







Bitcoin and Blockchain

Bitcoin: A Peer-to-Peer Electronic Cash System

Satoshi Nakamoto
satoshin@gmx.com
www.bitcoin.org

Abstract. A purely peer-to-peer version of electronic cash would allow online payments to be sent directly from one party to another without going through a financial institution. Digital signatures provide part of the solution, but the main benefits are lost if a trusted third party is still required to prevent double-spending. We propose a solution to the double-spending problem using a peer-to-peer network. The network timestamps transactions by hashing them into an ongoing chain of hash-based proof-of-work, forming a record that cannot be changed without redoing the proof-of-work. The longest chain not only serves as proof of the sequence of events witnessed, but proof that it came from the largest pool of CPU power. As long as a majority of CPU power is controlled by nodes that are not cooperating to attack the network, they'll generate the longest chain and outpace attackers. The network itself requires minimal structure. Messages are broadcast on a best effort basis, and nodes can leave and rejoin the network at will, accepting the longest proof-of-work chain as proof of what happened while they were gone.

1. Introduction

Commerce on the Internet has come to rely almost exclusively on financial institutions serving as trusted third parties to process electronic payments. While the system works well enough for





Bitcoin and Blockchain

Cheats on the CAP theorem

Eventual consistency

Availability

Partition-tolerance (**But short**)



Proof of Work



Proof of Work

Pricing via Processing or Combatting Junk Mail

Cynthia Dwork and Moni Naor

IBM Almaden Research Center
650 Harry Road
San Jose, CA 95120

Abstract. We present a computational technique for pricing access to e-mail in particular and controlling access to a shared resource in general. The main idea is to require a user to compute a non-trivial, not intractable, function in order to gain access to the resource, preventing frivolous use. To this end we suggest several models based on, respectively, extracting square roots modulo a prime, the Shamir signature scheme, and the Ong-Schnorr-Shamir signature scheme.



1 Introduction

Recently, one of us returned from a brief vacation, only to find 241 messages in our reader. While junk mail has long been a nuisance in hard (snail) mail, we

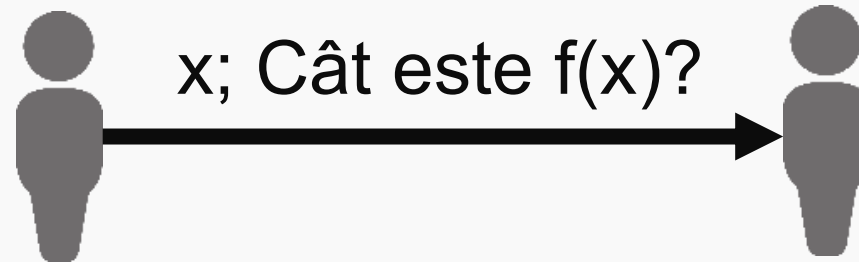


Proof of Work

- $y = f(x)$
 - $f()$ este extrem de lentă computațional
- $x = g(y)$
 - $g()$ este extrem de rapidă computațional



Proof of Work





Proof of Work



Processing!!!




Proof of Work



Processing!!!



Proof of Work





Proof of Work




$g(y) == x?$

DA - știu că a muncit

NU – încearcă să trișeze



Blockchain





Bitcoin energy consumption

Bitcoin Mining and its Energy Footprint

Karl J. O'Dwyer[†] and David Malone^{*}

*Hamilton Institute
National University of Ireland Maynooth*

E-mail: [†]`karl.a.odwyer@nuim.ie`

^{*}`david.malone@nuim.ie`

Abstract — Bitcoin is a digital cryptocurrency that has generated considerable public interest, including both booms in value and busts of exchanges dealing in Bitcoins. One of the fundamental concepts of Bitcoin is that work, called mining, must be done in checking all monetary transactions, which in turn creates Bitcoins as a reward. In this paper we look at the energy consumption of Bitcoin mining. We consider if and when Bitcoin mining has been profitable compared to the energy cost of performing the mining, and conclude that specialist hardware is usually required to make Bitcoin mining profitable. We also show that the power currently used for Bitcoin mining is comparable to Ireland's electricity consumption.



Bitcoin energy consumption




Fig. 5: Estimated Power Consumption of the Bitcoin Mining Network.





Transactions



ACID

Principles of Transaction-Oriented Database Recovery

THEO HAERDER


Fachbereich Informatik, University of Kaiserslautern, West Germany

ANDREAS REUTER¹

IBM Research Laboratory, San Jose, California 95193



er, a terminological framework is provided for descri
covery schemes for database systems in a conceptual
ation-dependent way. By introducing the terms mate
n strategy, and checkpoint, we obtain a means for cla
ations from a unified viewpoint. This is complement
logging techniques, which are precisely defined by u
these criteria are related to all relevant questions of





ACID

- Atomicity
- Consistency
- Isolation
- Durability





Distributed Hash Table Chord

Chord: A Scalable Peer-to-peer Lookup Service for Internet Applications

Ion Stoica; Robert Morris, David Karger, M. Frans Kaas
MIT Laboratory for Computer Science
chord@lcs.mit.edu
<http://pdos.lcs.mit.edu/chord/>

Abstract

A fundamental problem that confronts peer-to-peer applications is to efficiently locate the node that stores a particular data item. This paper presents *Chord*, a distributed lookup protocol that addresses this problem. Chord provides support for just one operation: given a key, it maps the key onto a node. Data location can be easily implemented on top of Chord by associating a key with each data item, and storing the key/data item pair at the node to which the key maps. Chord adapts efficiently as nodes join and leave the system, and can answer queries even if the system is continuously changing. Results from theoretical analysis, simulations, and experiments show that Chord is scalable, with communication cost and the state maintained by each node scaling logarithmically with the number of Chord nodes.

and involves relative
and leave the system

Previous work on
aware of most other
scale to large number
“routing” information
routing table is distributed
communicating with
an N -node system,
 $O(\log N)$ other nodes
messages to other nodes
nodes join and leave
event results in no re

Three features of the
peer lookup protocol





Amazon DynamoDB



Peer to Peer